

Accelerating Data Flow with Kafka: Use Cases, Advantages, and Why It's Lightning-Fast

What is Kafka?

Kafka is a decentralized streaming platform developed by LinkedIn in 2011 as a high-throughput message broker for its own use, then open sourced to the Apache Software Foundation.

It is designed to handle real-time data feeds and provides a high-throughput, resilient and scalable solution for processing and storing streams of records.

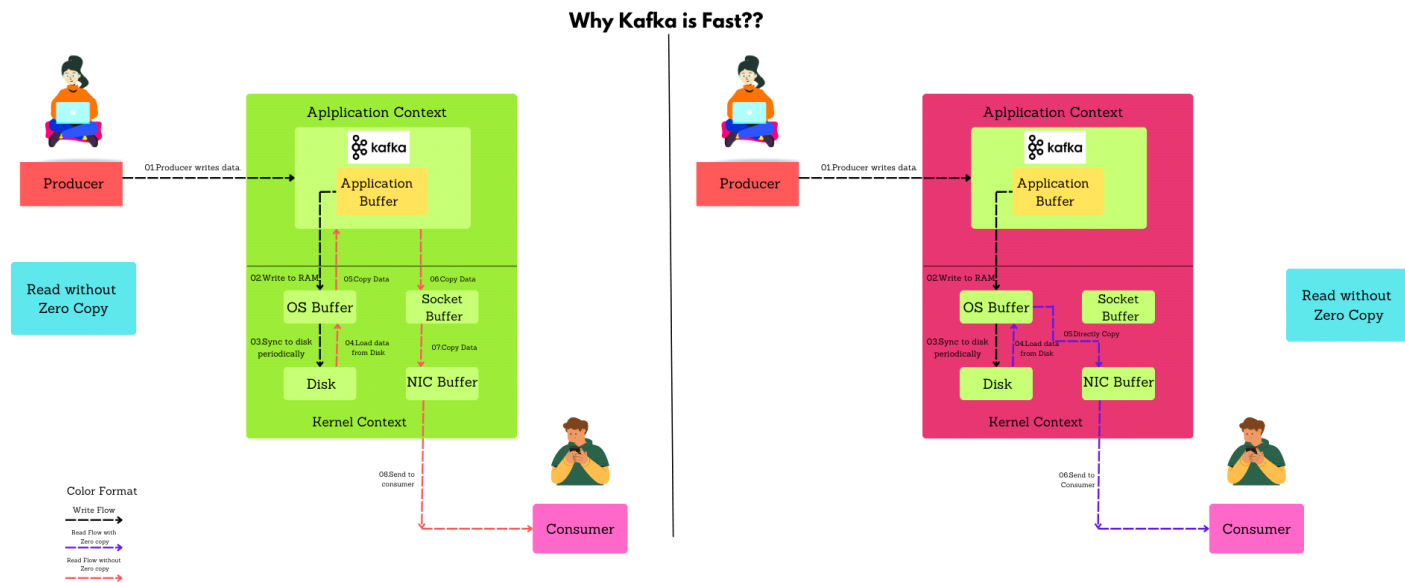
It is built on a publish-subscribe model where producers write data to topics and consumers read data from topics.

Kafka is often used in use cases such as log aggregation, stream processing, event sourcing, and real-time analytics.

Kafka provides durability and resilience by replicating data across multiple brokers in a cluster.

Kafka is an open-source distributed event streaming platform, which means it is free to use. However, there are some paid versions and services available from companies like Confluent that offer additional features and support for Kafka as per uses we need to purchase it.

How is Kafka so fast?



Kafka's exceptional speed is coordinated by two key virtuosos:

Sequential I/O and the Zero Copy Principle.

1.Sequential I/O -

Kafka addresses the perceived slowness of disks by implementing Sequential I/O in a brilliant manner. By utilizing a log structure for message storage, Kafka achieves speeds that are close to RAM, optimizing disk access and greatly reducing seek times. The outcome? Fast and low-latency message delivery.

2. Zero Copy Principle -

In the world of data transfer, Kafka stands out with its implementation of the Zero Copy Principle. By avoiding unnecessary data copies and minimizing context switches between user and kernel modes, Kafka significantly improves efficiency. The key concept behind this principle is requesting the kernel to directly move data to the NIC (Network Interface Controller) buffer. As a result, transfer time is remarkably reduced.

Fine-Tuned Efficiency:

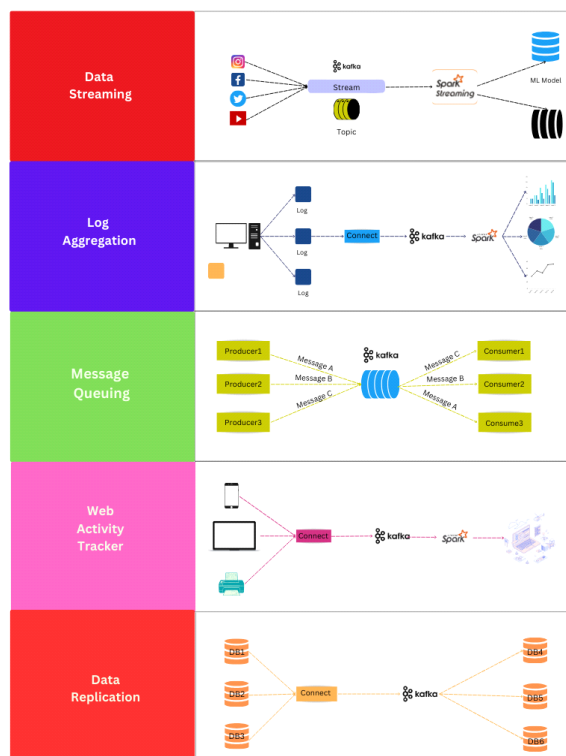
Beyond these titans, Kafka employs additional techniques to boost its prowess:

- **Message Compression:** Shrinking message sizes for faster transmission is a game-changer when it comes to handling substantial data volumes.
- **Message Batching:** Bundling messages together to streamline processing is particularly beneficial, especially in high-volume message environments.

In essence, Kafka's lightning speed is achieved through a combination of optimized architecture, smart compression techniques, batch processing capabilities, and efficient utilization of in-memory storage.

Use cases of Kafka

Top 5 Kafka Use cases



1. Data Streaming: Kafka is a powerful platform for processing and analyzing real-time data streams. Its streaming capabilities are unmatched, making it an excellent choice for building dynamic and responsive applications. Kafka integrates seamlessly with stream processing frameworks such as Kafka Streams, Apache Spark, and Apache Flink. It enables real-time processing of streaming data and provides support for complex event processing, transformations, and analytics.

2. Log Aggregation: Simplify your log management with Kafka! Centralize logs from various sources, making troubleshooting and analysis a breeze. Say goodbye to scattered logs and hello to a unified log aggregation solution.

Kafka is commonly used for log aggregation, where it collects log data from various sources and centralizes them into a single location. This helps in easy monitoring, analysis, and troubleshooting of distributed systems.

3. Messaging Queue: Kafka is an excellent messaging system that excels in high-throughput and fault-tolerance. It enhances communication between microservices, applications, and components with its reliable message queue architecture. One of the most important uses of Kafka is as a distributed messaging system. It is commonly used for building real-time data pipelines and streaming applications. Kafka allows for fast, efficient, fault-tolerant, and scalable message processing.

4. Web Activity Tracker: Stay updated on user interactions and behaviors in real-time. Kafka's capability to handle high volumes of data and offer low-latency processing makes it a perfect option for developing reliable web activity tracking systems.

5. Data Replication: Kafka's data replication capabilities ensure data consistency across distributed systems. Whether you need to replicate data between data centers or ensure high availability, Kafka has you covered. You can also use Kafka for data replication and backup purposes, allowing you to replicate data across multiple data centers or cloud regions and ensuring data availability and disaster recovery.

6) Data Integration: Kafka provides reliable and scalable data integration between different systems and applications. It enables seamless data transfer between various databases, data lakes, and data warehouses. Companies can extract value from dispersed data by connecting with diverse data sources and destinations.

7) Event Streaming: Kafka is commonly used for event streaming, serving as a central hub for capturing, storing, and processing events in real-time. It enables the implementation of event-driven architectures and simplifies the integration of different systems.

8) Internet of Things (IoT): Kafka is commonly used in IoT applications to collect and process real-time data from a variety of sensors and devices. It offers a scalable and dependable platform for managing the significant amounts of streaming data generated by IoT devices.

9) Commit Log: Kafka stores all the messages in a distributed commit log, which ensures durability and fault-tolerance. This makes it suitable for applications that require reliable message delivery and storage.

10) Microservices Communication: Kafka can be used as a communication channel between microservices in a distributed system. It enables the decoupling of services and provides a scalable and fault-tolerant approach to exchanging messages and events.

Kafka Programs with installation and dotnet application with steps

Here, I will provide a basic example of how to install Kafka on your machine and write programs in .NET Core 6.0 to connect with a Kafka application for message queueing.

In this scenario, I create two applications: one is the Producer, and the second is the Consumer. The Producers write messages and transfer them to the Consumer via Kafka.

Here, I will provide detailed steps for installing Kafka along with example programs.

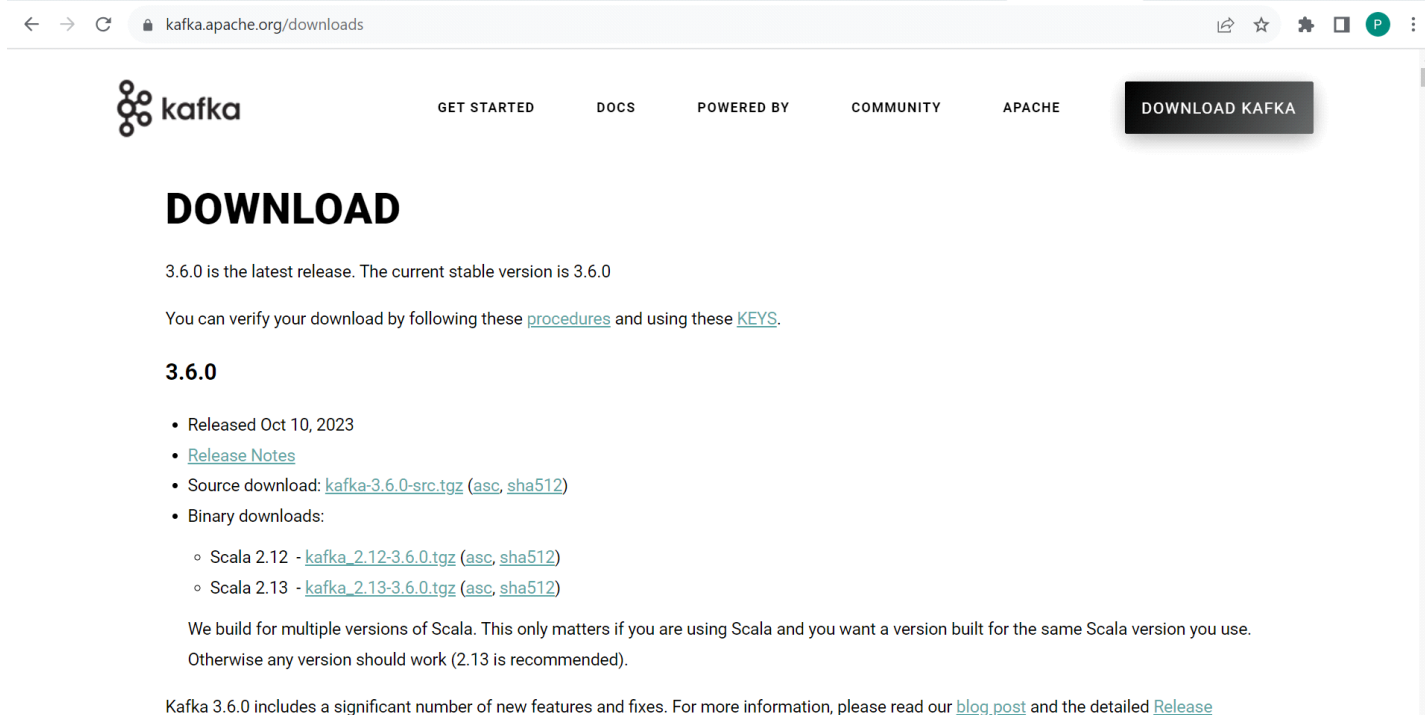
How to install Kafka in your machine .

1) How to download Kafka:

Here I am providing you a link. You can download from this link.

Link - <https://kafka.apache.org/downloads>

You can download the Kafka version that meets your requirements. Here is the latest version of Kafka

The screenshot shows the 'Downloads' page of the Apache Kafka website. The browser's address bar displays 'kafka.apache.org/downloads'. The website's navigation bar includes the Kafka logo, links for 'GET STARTED', 'DOCS', 'POWERED BY', 'COMMUNITY', and 'APACHE', and a prominent 'DOWNLOAD KAFKA' button. The main heading is 'DOWNLOAD'. Below it, a paragraph states that 3.6.0 is the latest release and the current stable version. It provides instructions on how to verify the download using 'procedures' and 'KEYS'. A section for version '3.6.0' lists release details: it was released on Oct 10, 2023, and provides links to 'Release Notes', the source download ('kafka-3.6.0-src.tgz'), and binary downloads for Scala 2.12 and 2.13. A note mentions that the binaries are built for multiple versions of Scala, with 2.13 being recommended. At the bottom, it mentions that Kafka 3.6.0 includes new features and fixes, with links to a 'blog.post' and a 'Release' page.

← → ↻ 🔒 kafka.apache.org/downloads

kafka GET STARTED DOCS POWERED BY COMMUNITY APACHE **DOWNLOAD KAFKA**

DOWNLOAD

3.6.0 is the latest release. The current stable version is 3.6.0

You can verify your download by following these [procedures](#) and using these [KEYS](#).

3.6.0

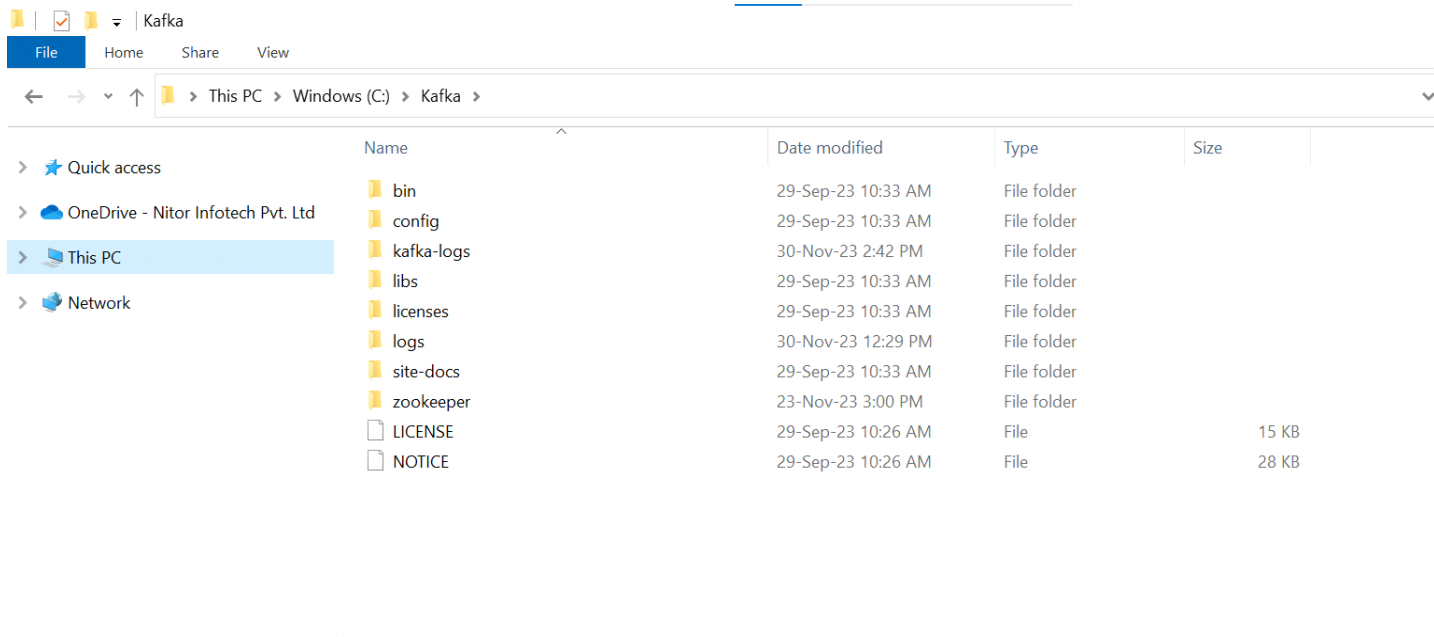
- Released Oct 10, 2023
- [Release Notes](#)
- Source download: [kafka-3.6.0-src.tgz](#) ([asc](#), [sha512](#))
- Binary downloads:
 - Scala 2.12 - [kafka_2.12-3.6.0.tgz](#) ([asc](#), [sha512](#))
 - Scala 2.13 - [kafka_2.13-3.6.0.tgz](#) ([asc](#), [sha512](#))

We build for multiple versions of Scala. This only matters if you are using Scala and you want a version built for the same Scala version you use. Otherwise any version should work (2.13 is recommended).

Kafka 3.6.0 includes a significant number of new features and fixes. For more information, please read our [blog.post](#) and the detailed [Release](#)

After downloading, you can create a folder named 'Kafka' in the C: drive and paste the downloaded file into that folder.

After extracting and pasting, the folder will look like this.



We have two important files that need to be executed when running Kafka: the server.properties file and the zookeeper.properties file, both located inside the config folder.

Please ensure that the JDK is already installed on your machine.

Open two terminals with same path Kafka -> " C:\Kafka> "

2) How to run Kafka server and Zookeeper.

And hit the following commands to run Zookeeper and the server.

Server -> .\bin\windows\kafka-server-start.bat config\server.properties

```
C:\Windows\System32\cmd.exe - .\bin\windows\kafka-server-start.bat config\server.properties
Microsoft Windows [Version 10.0.19045.3570]
(c) Microsoft Corporation. All rights reserved.

C:\Kafka>.bin\windows\kafka-server-start.bat config\server.properties
[2023-11-30 12:29:48,639] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log4jControllerRegistration$)
[2023-11-30 12:29:48,981] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disable client-initiated TLS renegotiation (org.apache.zookeeper.common.X509Util
1)
[2023-11-30 12:29:49,050] INFO starting (kafka.server.KafkaServer)
[2023-11-30 12:29:49,051] INFO Connecting to zookeeper on localhost:2181 (kafka.server.KafkaServer)
[2023-11-30 12:29:49,065] INFO [ZooKeeperClient Kafka server] Initializing a new session to localhost:2181. (kafka.zookeeper.ZooKeeperClient)
[2023-11-30 12:29:49,071] INFO Client environment:zookeeper.version=3.8.2-139d619b58292d7734b4fc83a0f44be4e7b0c986, built on 2023-07-05 19:24 UTC (org.apache.zookeeper.Zook
eeper)
[2023-11-30 12:29:49,071] INFO Client environment:host.name=NIT-LPT-R855.nitorinfotech.net (org.apache.zookeeper.ZooKeeper)
[2023-11-30 12:29:49,072] INFO Client environment:java.version=1.8.0_391 (org.apache.zookeeper.ZooKeeper)
[2023-11-30 12:29:49,072] INFO Client environment:java.vendor=Oracle Corporation (org.apache.zookeeper.ZooKeeper)
[2023-11-30 12:29:49,072] INFO Client environment:java.home=C:\Program Files\Java\jre-1.8 (org.apache.zookeeper.ZooKeeper)
[2023-11-30 12:29:49,072] INFO Client environment:java.class.path=C:\Kafka\libs\activation-repackaged-2.6.1.jar;C:\Kafka\libs\argparse4j
-0.7.0.jar;C:\Kafka\libs\audience-annotations-0.12.0.jar;C:\Kafka\libs\caffeine-2.9.3.jar;C:\Kafka\libs\checker-qual-3.19.0.jar;C:\Kafka\libs\commons-beanutils-1.9.4.jar;C:\
Kafka\libs\commons-cli-1.4.jar;C:\Kafka\libs\commons-collections-3.2.2.jar;C:\Kafka\libs\commons-digester-2.1.jar;C:\Kafka\libs\commons-io-2.11.0.jar;C:\Kafka\libs\commons
-lang3-3.8.1.jar;C:\Kafka\libs\commons-logging-1.2.jar;C:\Kafka\libs\commons-validator-1.7.jar;C:\Kafka\libs\connect-api-3.6.0.jar;C:\Kafka\libs\connect-basic-auth-extensio
n-3.6.0.jar;C:\Kafka\libs\connect-file-3.6.0.jar;C:\Kafka\libs\connect-json-3.6.0.jar;C:\Kafka\libs\connect-mirror-3.6.0.jar;C:\Kafka\libs\connect-mirror-client-3.6.0.jar;C
:\Kafka\libs\connect-runtime-3.6.0.jar;C:\Kafka\libs\connect-transforms-3.6.0.jar;C:\Kafka\libs\error-prone-annotations-2.10.0.jar;C:\Kafka\libs\hk2-api-2.6.1.jar;C:\Kafka\
libs\hk2-locator-2.6.1.jar;C:\Kafka\libs\hk2-utils-2.6.1.jar;C:\Kafka\libs\jackson-annotations-2.13.5.jar;C:\Kafka\libs\jackson-core-2.13.5.jar;C:\Kafka\libs\jackson-databi
nd-2.13.5.jar;C:\Kafka\libs\jackson-dataformat-csv-2.13.5.jar;C:\Kafka\libs\jackson-datatype-jdk8-2.13.5.jar;C:\Kafka\libs\jackson-jaxrs-base-2.13.5.jar;C:\Kafka\libs\jacks
on-jaxrs-json-provider-2.13.5.jar;C:\Kafka\libs\jackson-module-jaxb-annotations-2.13.5.jar;C:\Kafka\libs\jackson-module-scala_2.13-2.13.5.jar;C:\Kafka\libs\jakarta.activati
on-api-1.2.2.jar;C:\Kafka\libs\jakarta.annotation-api-1.3.5.jar;C:\Kafka\libs\jakarta.inject-2.6.1.jar;C:\Kafka\libs\jakarta.validation-api-2.0.2.jar;C:\Kafka\libs\jakarta.
ws-rs-api-2.1.6.jar;C:\Kafka\libs\jakarta.xml.bind-api-2.3.3.jar;C:\Kafka\libs\javassist-3.29.2-GA.jar;C:\Kafka\libs\javax.activation-api-1.2.0.jar;C:\Kafka\libs\javax.anno
tation-api-1.3.2.jar;C:\Kafka\libs\javax.servlet-api-3.1.0.jar;C:\Kafka\libs\javax.ws.rs-api-2.1.1.jar;C:\Kafka\libs\jaxb-api-2.3.1.jar;C:\Kafka\libs\jersey-client-2.39.1.j
ar;C:\Kafka\libs\jersey-common-2.39.1.jar;C:\Kafka\libs\jersey-container-servlet-2.39.1.jar;C:\Kafka\libs\jersey-container-servlet-core-2.39.1.jar;C:\Kafka\libs\jersey-hk2-2.39.1.j
ar;C:\Kafka\libs\jersey-server-2.39.1.jar;C:\Kafka\libs\jetty-client-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-continuation-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-
http-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-io-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-security-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-server-9.4.52.v20230823.jar;C:\K
afka\libs\jetty-servlet-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-servlets-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-util-9.4.52.v20230823.jar;C:\Kafka\libs\jetty-util-ajax-9
.4.52.v20230823.jar;C:\Kafka\libs\jline-3.22.0.jar;C:\Kafka\libs\jopt-simple-5.0.4.jar;C:\Kafka\libs\jose4j-0.9.3.jar;C:\Kafka\libs\jsr305-3.0.2.jar;C:\Kafka\libs\kafka-cli
ents-3.6.0.jar;C:\Kafka\libs\kafka-group-coordinator-3.6.0.jar;C:\Kafka\libs\kafka-log4j-appender-3.6.0.jar;C:\Kafka\libs\kafka-metadata-3.6.0.jar;C:\Kafka\libs\kafka-raft-3.6.0.j
ar;C:\Kafka\libs\kafka-server-common-3.6.0.jar;C:\Kafka\libs\kafka-shell-3.6.0.jar;C:\Kafka\libs\kafka-storage-3.6.0.jar;C:\Kafka\libs\kafka-storage-api-3.6.0.jar;C:\
Kafka\libs\kafka-streams-3.6.0.jar;C:\Kafka\libs\kafka-streams-examples-3.6.0.jar;C:\Kafka\libs\kafka-streams-scala_2.13-3.6.0.jar;C:\Kafka\libs\kafka-streams-test-utils-3
.6.0.jar;C:\Kafka\libs\kafka-tools-3.6.0.jar;C:\Kafka\libs\kafka-tools-api-3.6.0.jar;C:\Kafka\libs\kafka_2.13-3.6.0.jar;C:\Kafka\libs\lz4-java-1.8.0.jar;C:\Kafka\libs\maven
-artifact-3.8.8.jar;C:\Kafka\libs\metrics-core-2.2.0.jar;C:\Kafka\libs\metrics-core-4.1.12.1.jar;C:\Kafka\libs\netty-buffer-4.1.94.Final.jar;C:\Kafka\libs\netty-codec-4.1.9
4.Final.jar;C:\Kafka\libs\netty-common-4.1.94.Final.jar;C:\Kafka\libs\netty-handler-4.1.94.Final.jar;C:\Kafka\libs\netty-resolver-4.1.94.Final.jar;C:\Kafka\libs\netty-trans
port-4.1.94.Final.jar;C:\Kafka\libs\netty-transport-classes-epoll-4.1.94.Final.jar;C:\Kafka\libs\netty-transport-native-epoll-4.1.94.Final.jar;C:\Kafka\libs\netty-transport
-native-unix-common-4.1.94.Final.jar;C:\Kafka\libs\osgi-resource-locator-1.0.3.jar;C:\Kafka\libs\paranamer-2.8.jar;C:\Kafka\libs\pcollections-4.0.1.jar;C:\Kafka\libs\plexus
```

After executing this command, your server will start on Port 9092.

Use this port when connecting your .NET application to the Kafka server."

Zookeeper -> .\bin\windows\zookeeper-server-start.bat config\zookeeper.properties


```
C:\Windows\System32\cmd.exe - .\bin\windows\zookeeper-server-start.bat config\zookeeper.properties
Microsoft Windows [Version 10.0.19045.3570]
(c) Microsoft Corporation. All rights reserved.

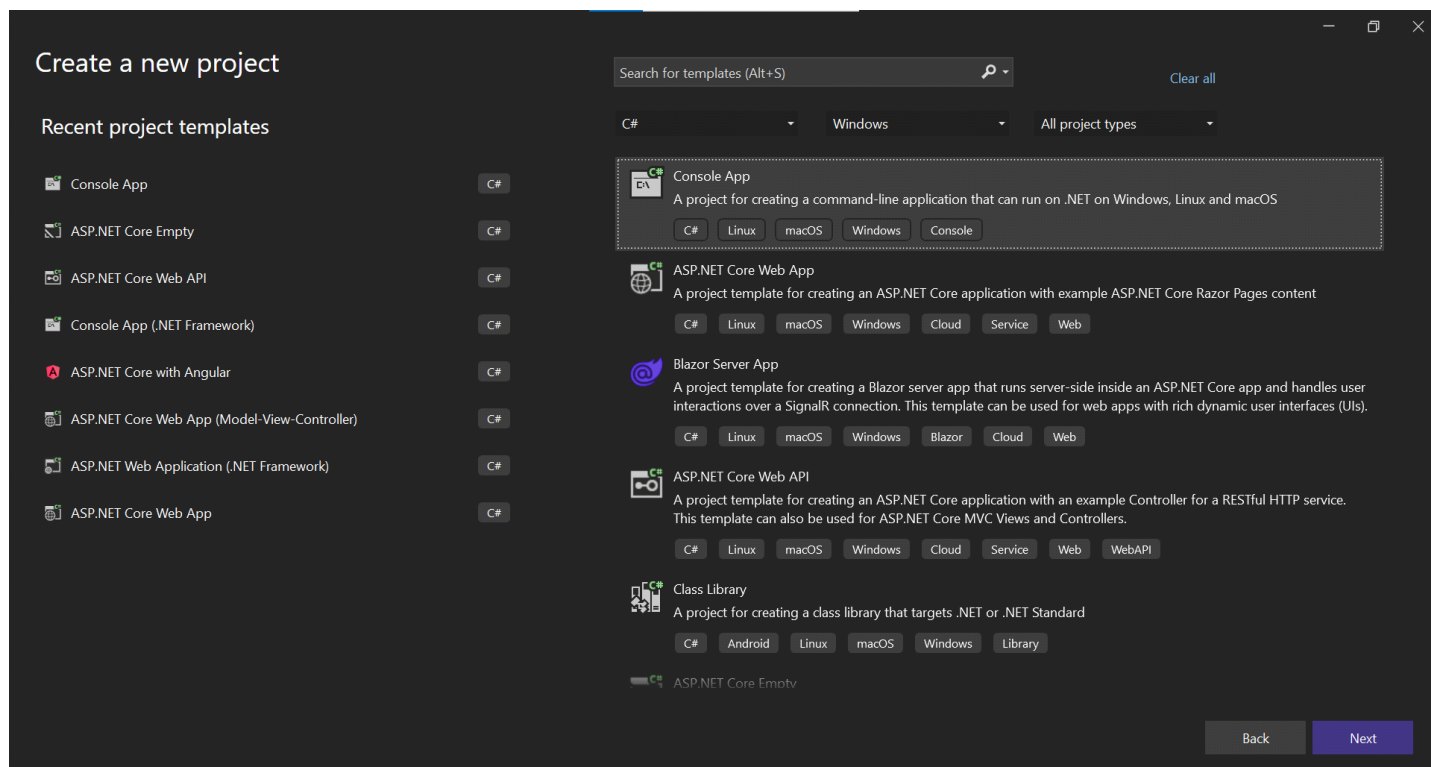
C:\Kafka>.bin\windows\zookeeper-server-start.bat config\zookeeper.properties
[2023-11-30 12:29:15,028] INFO Reading configuration from: config\zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,031] WARN config\zookeeper.properties is relative. Prepend .\ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,041] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,041] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,041] INFO observerMasterPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,042] INFO metricsProvider.className is org.apache.zookeeper.metrics.impl.DefaultMetricsProvider (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,043] INFO autopurge.snapRetainCount set to 3 (org.apache.zookeeper.server.DataDirCleanupManager)
[2023-11-30 12:29:15,044] INFO autopurge.purgeInterval set to 0 (org.apache.zookeeper.server.DataDirCleanupManager)
[2023-11-30 12:29:15,044] INFO Purge task is not scheduled. (org.apache.zookeeper.server.DataDirCleanupManager)
[2023-11-30 12:29:15,044] WARN Either no config or no quorum defined in config, running in standalone mode (org.apache.zookeeper.server.quorum.QuorumPeerMain)
[2023-11-30 12:29:15,045] INFO Log4j 1.2 jmx support not found; jmx disabled. (org.apache.zookeeper.jmx.ManagedUtil)
[2023-11-30 12:29:15,046] INFO Reading configuration from: config\zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,046] WARN config\zookeeper.properties is relative. Prepend .\ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,048] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,049] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,049] INFO observerMasterPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,049] INFO metricsProvider.className is org.apache.zookeeper.metrics.impl.DefaultMetricsProvider (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2023-11-30 12:29:15,049] INFO Starting server (org.apache.zookeeper.server.ZooKeeperServerMain)
[2023-11-30 12:29:15,061] INFO ServerMetrics initialized with provider org.apache.zookeeper.metrics.impl.DefaultMetricsProvider@90f6bfd (org.apache.zookeeper.server.ServerMetrics)
[2023-11-30 12:29:15,064] INFO ACL digest algorithm is: SHA1 (org.apache.zookeeper.server.auth.DigestAuthenticationProvider)
[2023-11-30 12:29:15,064] INFO zookeeper.DigestAuthenticationProvider.enabled = true (org.apache.zookeeper.server.auth.DigestAuthenticationProvider)
[2023-11-30 12:29:15,073] INFO zookeeper.snapshot.trust.empty : false (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2023-11-30 12:29:15,083] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,084] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,084] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,086] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,087] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,087] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,088] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,088] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,088] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,089] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,089] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,089] INFO (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,095] INFO Server environment:zookeeper.version=3.8.2-139d619b58292d7734b4fc83a0f44be4e7b0c986, built on 2023-07-05 19:24 UTC (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,095] INFO Server environment:host.name=NIT-LPT-R855.nitorinfotech.net (org.apache.zookeeper.server.ZooKeeperServer)
[2023-11-30 12:29:15,096] INFO Server environment:java.version=1.8.0_391 (org.apache.zookeeper.server.ZooKeeperServer)
```

After executing these two commands in separate terminals, you can start your Visual Studio instance.

You need to create two applications for consuming and producing Kafka messages: one for the Consumer and one for the Producer.

3) Create a new Project for kafka consumer -

Select console application in .Net 6 C#



You can name it "KafkaConsumer." The application will be responsible for consuming Kafka messages.

Configure your new project

Console App

C#

Linux

macOS

Windows

Console

Project name

KafkaConsumer

Location

C:\Users\pavan.nangare\OneDrive - Nitor Infotech Pvt. Ltd\Desktop\BLOG\Kafka\K app\

...

Solution name ⓘ

KafkaConsumer

☐ Place solution and project in the same directory

Then, you can select .NET 6.0 (Long-Term Support) and click on Create

Additional information

Console App


C#


Linux

macOS


Windows

Console

Framework 

.NET 6.0 (Long Term Support) 

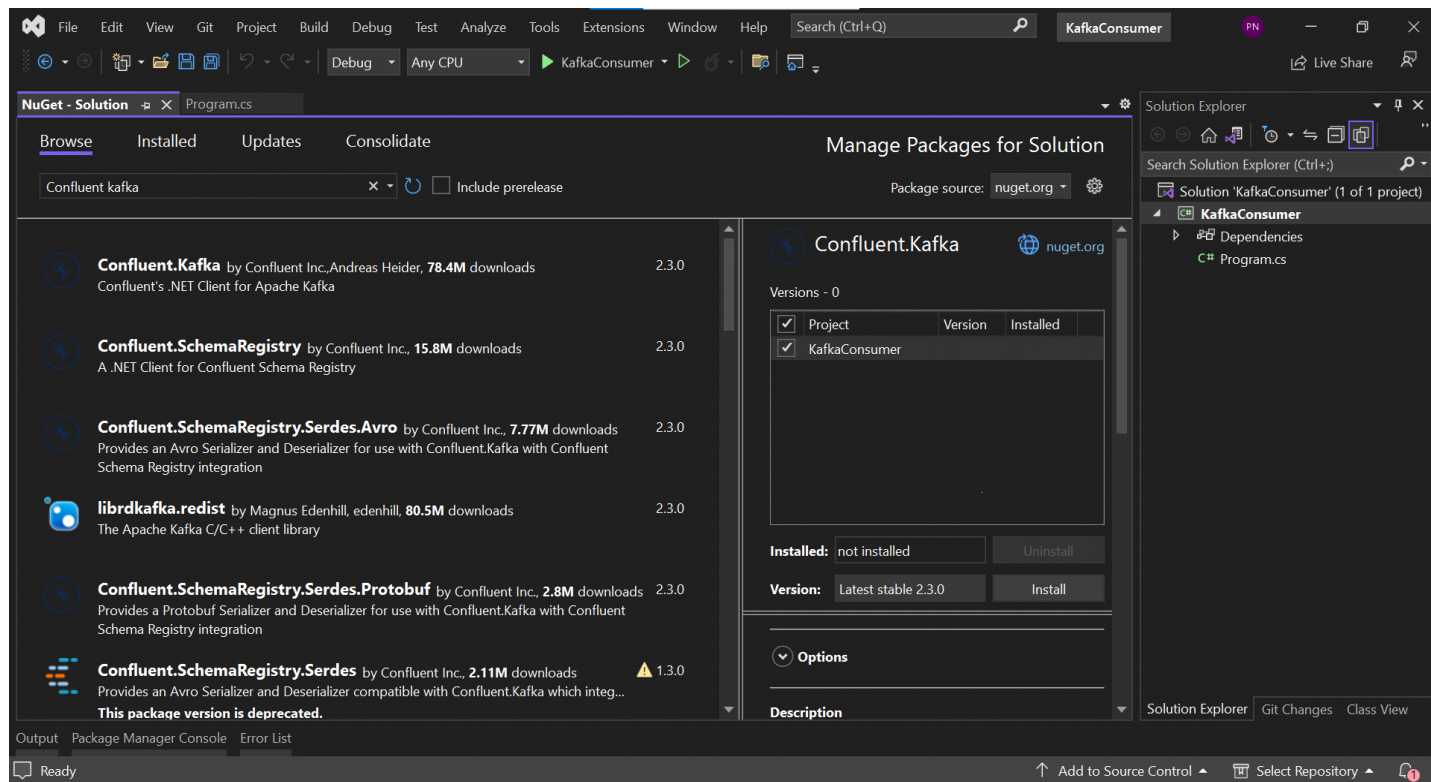
☐

Do not use top-level statements 

The application will be created successfully.

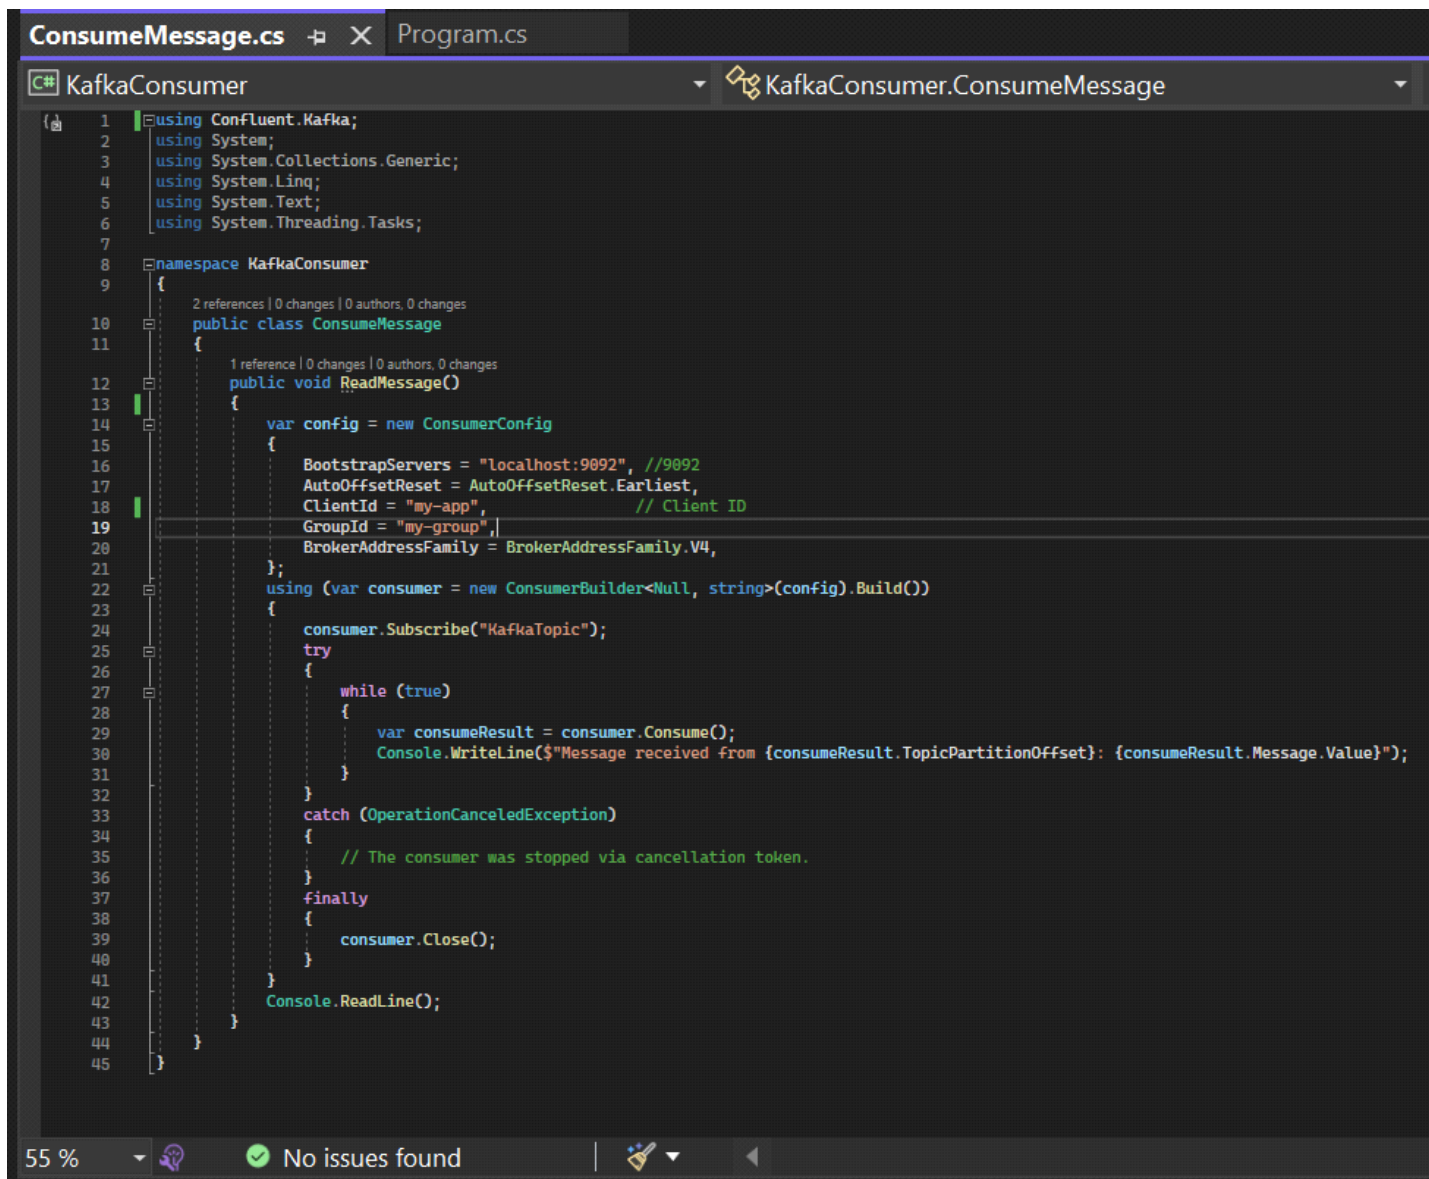
4) Install Confluent.Kafka package -

Once the application is started, you should add the package as "Confluent.Kafka"



5) Write consumer program-

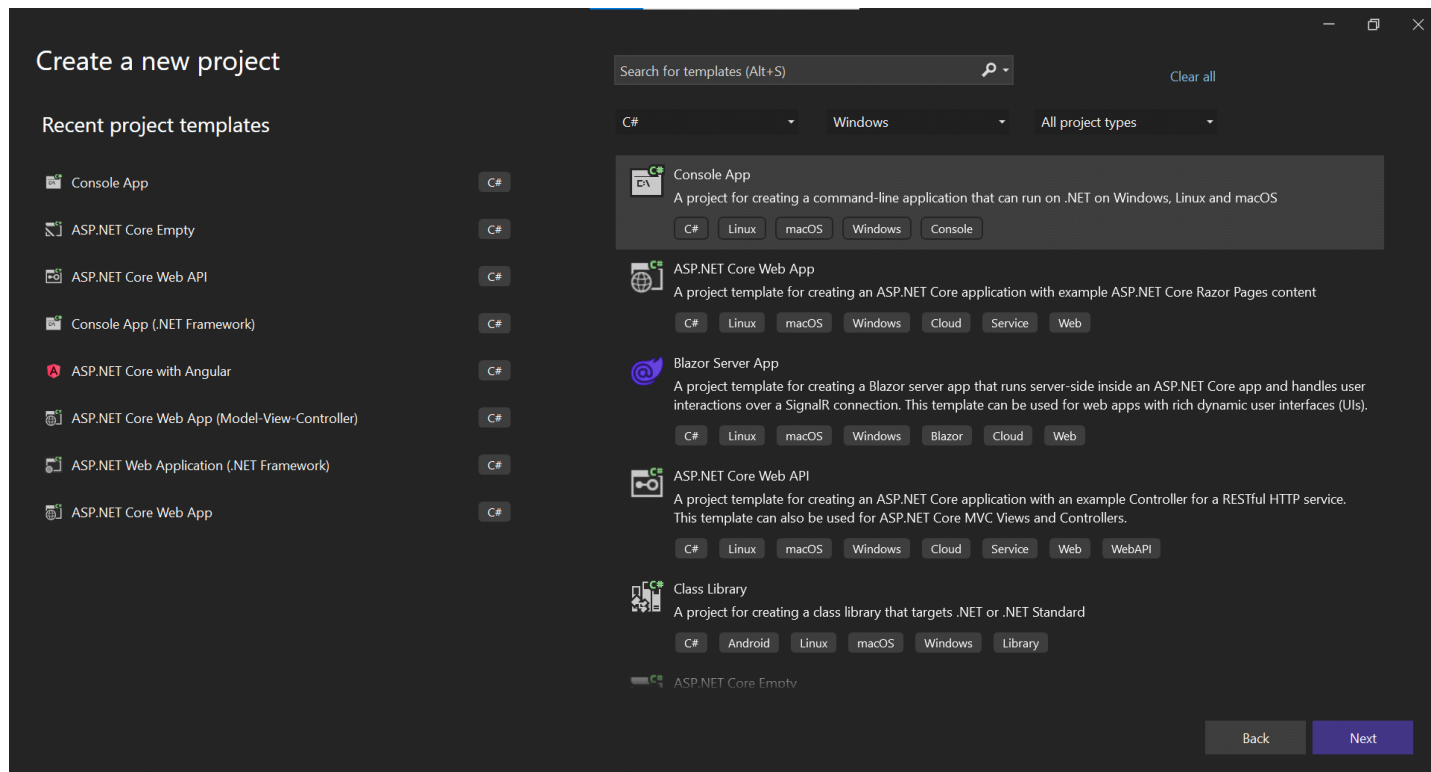
Write the code below for a consumer instance and create another instance for a producer.



```
1 using Confluent.Kafka;
2 using System;
3 using System.Collections.Generic;
4 using System.Linq;
5 using System.Text;
6 using System.Threading.Tasks;
7
8 namespace KafkaConsumer
9 {
10     2 references | 0 changes | 0 authors, 0 changes
11     public class ConsumeMessage
12     {
13         1 reference | 0 changes | 0 authors, 0 changes
14         public void ReadMessage()
15         {
16             var config = new ConsumerConfig
17             {
18                 BootstrapServers = "localhost:9092", //9092
19                 AutoOffsetReset = AutoOffsetReset.Earliest,
20                 ClientId = "my-app", // Client ID
21                 GroupId = "my-group",
22                 BrokerAddressFamily = BrokerAddressFamily.V4,
23             };
24             using (var consumer = new ConsumerBuilder<Null, string>(config).Build())
25             {
26                 consumer.Subscribe("KafkaTopic");
27                 try
28                 {
29                     while (true)
30                     {
31                         var consumeResult = consumer.Consume();
32                         Console.WriteLine($"Message received from {consumeResult.TopicPartitionOffset}: {consumeResult.Message.Value}");
33                     }
34                 }
35                 catch (OperationCanceledException)
36                 {
37                     // The consumer was stopped via cancellation token.
38                 }
39                 finally
40                 {
41                     consumer.Close();
42                 }
43             }
44             Console.ReadLine();
45         }
46     }
47 }
```

6) Create same application for kafka producer -

You can do the same for a Kafka Producer application. Select the console application for .NET 6.0 C#.



The second application is called "KafkaProducer," and its purpose is to produce Kafka messages.

Configure your new project

Console App

C#

Linux

macOS

Windows

Console

Project name

KafkaProducer

Location

C:\Users\pavan.nangare\OneDrive - Nitor Infotech Pvt. Ltd\Desktop\BLOG\Kafka\K app

...

Solution name [i](#)

KafkaProducer

☐

Place solution and project in the same directory

To keep the framework as .NET 6.0 (Long-term Support), click on the "Create" button.

Additional information

Console App


C#


Linux

macOS


Windows

Console

Framework 

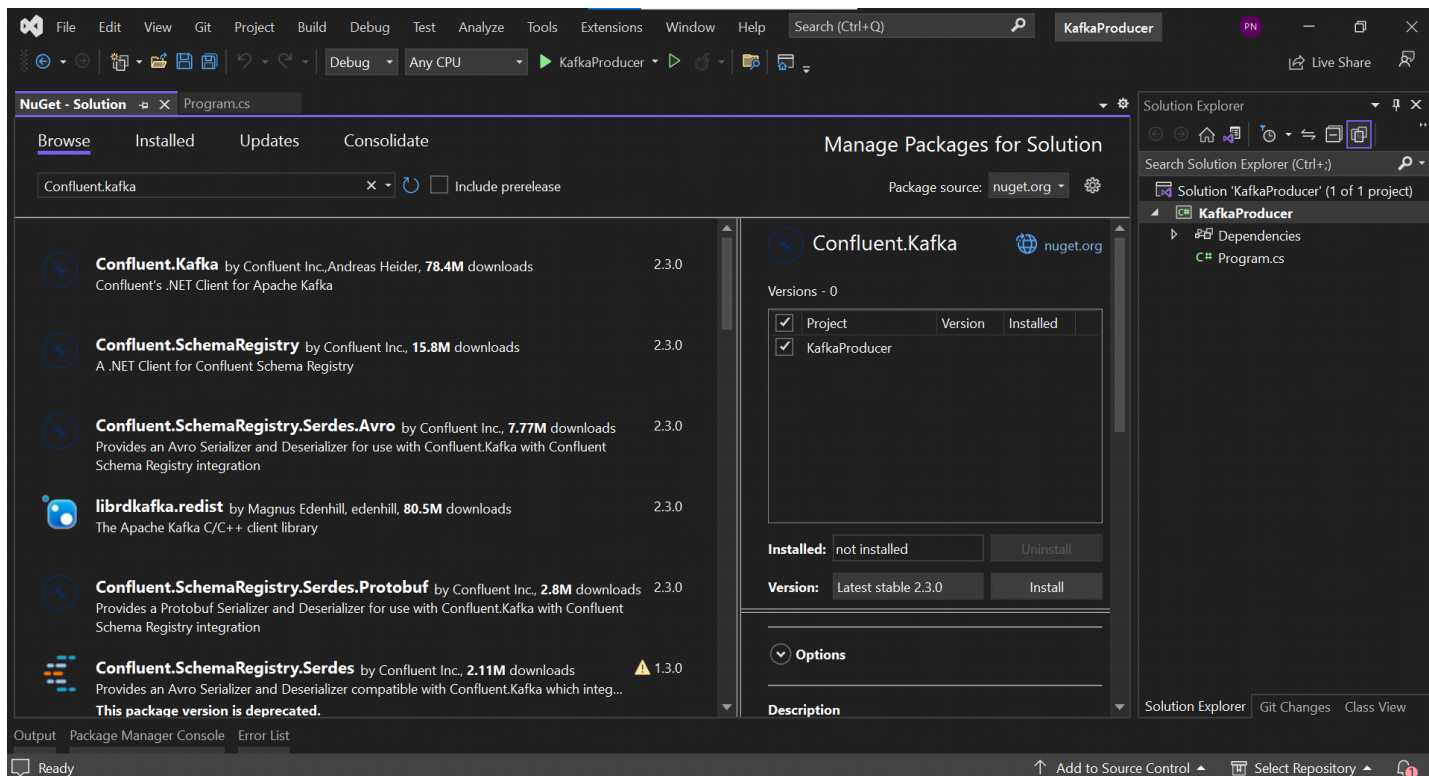
.NET 6.0 (Long Term Support) 

☐

Do not use top-level statements 

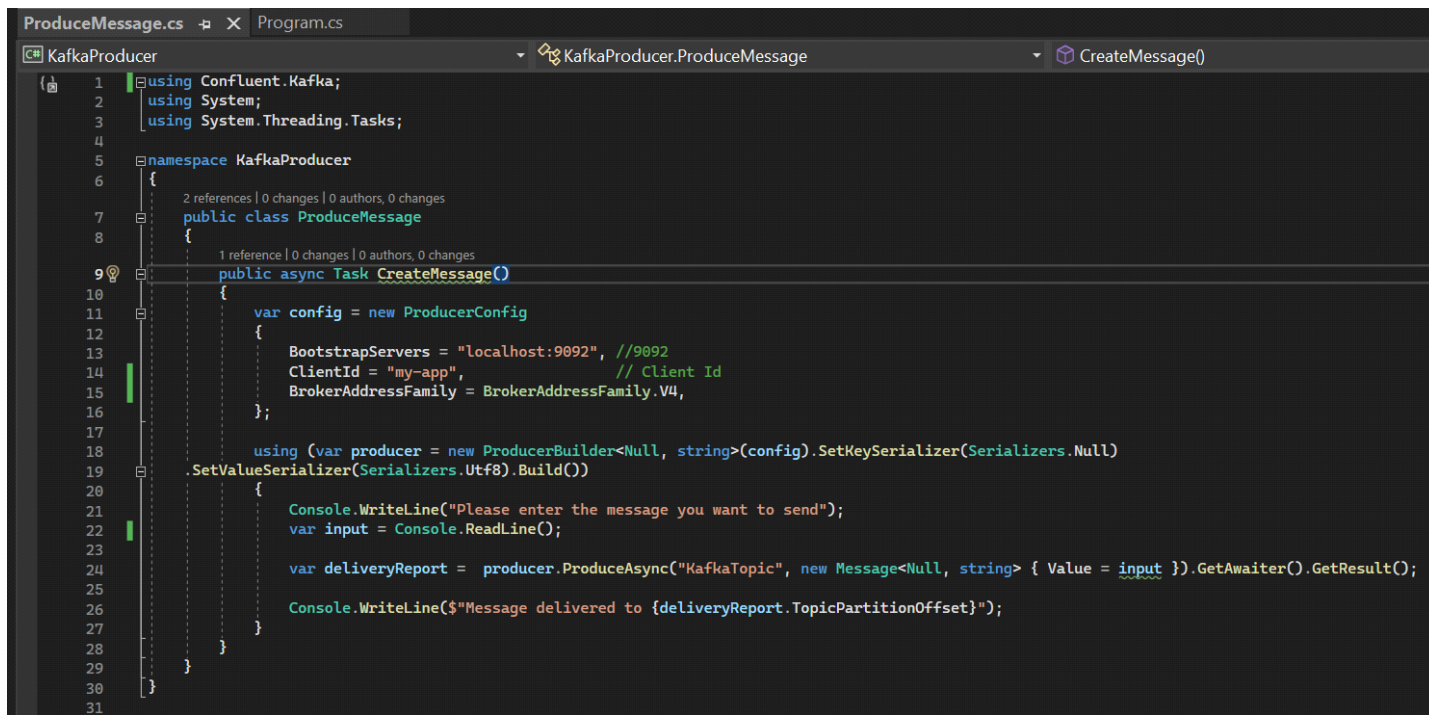
7) Install Confluent.Kafka Package -

After successfully creating a new Producer application, you should install the package as "Confluent.Kafka".



8) Write program for Kafka Producer

The program below is for Kafka producers.



9) Run both projects Kafka Consumer and Kafka Producer -

To run both the Consumer and Producer programs simultaneously, open two different terminals. One terminal will be for the Consumer program and the other for the Producer program.

10) Send message from kafka Producer

This is producer terminal

```
C:\Users\pavan.nangare\OneDrive - Nitor Infotech Pvt. Ltd\Desktop\BLOG\Kafka'
--Producer--
Please enter the message you want to send
Hello Nitorians..!!_
```

After hitting Enter, the message will be consumed on the consumer's terminal, and the window of the consumer's terminal is.

11) Receive messages sent from the Kafka producer-

```
C:\Users\pavan.nangare\OneDrive - Nitor Infotech Pvt. Ltd\Desktop\BLOG\Kafka\kafkaConsu
--Consumer--
Message received from KafkaTopic [[0]] @4: Hello Nitorians..!!
```

Why Kafka?

Kafka is a popular choice for building real-time streaming data pipelines and applications due to the following reasons:

- 1. High performance:** Kafka is renowned for its high throughput and low latency. It has the capability to handle millions of messages per second and is designed with a focus on writing data, making it well-suited for situations that necessitate rapid data ingestion and processing.
- 2. Persistent storage:** Kafka stores messages on disk, which enables durability and fault-tolerance. This means that even if a consumer goes offline, it can resume consuming messages from the point it left off once it comes back online.
- 3. Distributed architecture:** Kafka is designed as a distributed system, which means it can handle large amounts of data and scale horizontally by adding more machines to the cluster. This allows for high throughput and fault-tolerance.
- 4. Publish-subscribe model:** Kafka follows a publish-subscribe messaging pattern, where producers publish messages to topics, and consumers subscribe to those topics to receive the messages. This decoupling of producers and consumers enables easy scalability and flexibility in building data pipelines.
- 5. Fault-tolerance:** Kafka is designed to be highly available and fault-tolerant, achieving this by replicating data across multiple brokers in a cluster. In the event of a broker failure, data can still be accessed and consumed from other replicas.
- 6. Scalability:** Kafka's distributed architecture enables horizontal scalability. By adding additional brokers to the cluster, you can effectively handle higher data ingestion and processing demands while maintaining optimal performance.

Overall, Kafka's design principles and features make it an excellent choice for building fast, scalable, and fault-tolerant data pipelines and real-time streaming applications.

Git Link - <https://github.com/nangarepavan8/Kafka>